



AI on Linux on Z

Health insurance claims solution template

This solution template provides an example on how to deploy AI using a Linux on Z environment, while making use of open source frameworks, Triton Inference Server (TIS), and more.

Within this solution template, there are various phases of the AI lifecycle included. Work through each of the following steps to deploy your own health insurance claims solution on Linux on Z.



Table of contents

AI model training.....3

AI model deployment.....8

AI model integration.....12



Step 1

AI model training

We will build a health insurance claims AI model by training with the provided Rapid AI on Linux on Z Development Jupyter notebook. Simply point the Jupyter notebook to your dataset and run it to generate your AI model. This trained AI model can then be deployed with TIS.

All sample code for this section is within

```
aionz-st-health-insurance-claims-tis/zST-model-  
training-jupyter
```

Prerequisites

- Must have Python (3.9 or 3.10) installed

Dataset guidance

Sample health insurance claims dataset can be found on Kaggle -

https://www.kaggle.com/datasets/grvmishra7/health-insurance-prediction?select=train_Df64byy.csv

Required features

- applicant_id
- years_of_insurance_with_us
- regular_checkup_last_year
- adventure_sports
- occupation
- visited_doctor_last_1_year
- cholesterol_level
- daily_avg_steps
- age
- heart_decs_history
- any_other_major_decs_history gender
- avg_glucose_level
- bmi
- smoking_status
- year_last_admitted
- location
- alcohol
- exercise
- weight_change_in_last_one_year
- fat_percentage
- insurance_cost

Access rapid AI on Linux
on Z development
environment

Provide data

Model training

Access trained AI model

Access rapid AI on Linux on Z development environment

1. Access sample code

```
cd zST-model-training-jupyter
```

2. Create and activate Python virtual environment

```
python -m venv env
source env/bin/activate
```

3. Install required Python packages

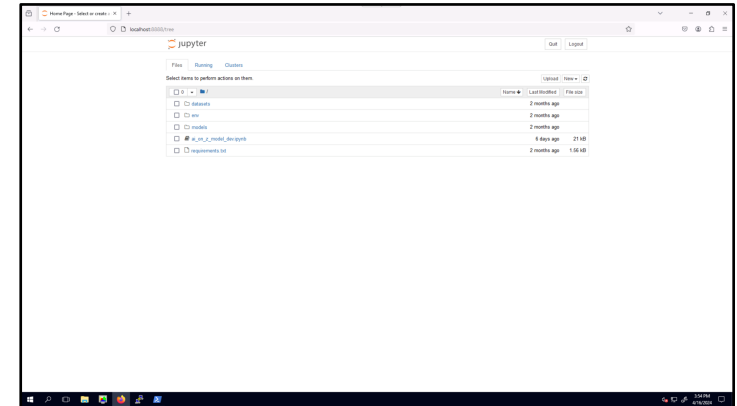
```
pip install -r requirements.txt
```

4. Run Jupyter

```
jupyter notebook
```

5. View Jupyter interface

- a. Go to localhost:8888 in a web browser



1. AI model training

Access rapid AI on Linux
on Z development
environment

Provide data

Model training

Access trained AI model

2. AI model deployment

3. AI model integration

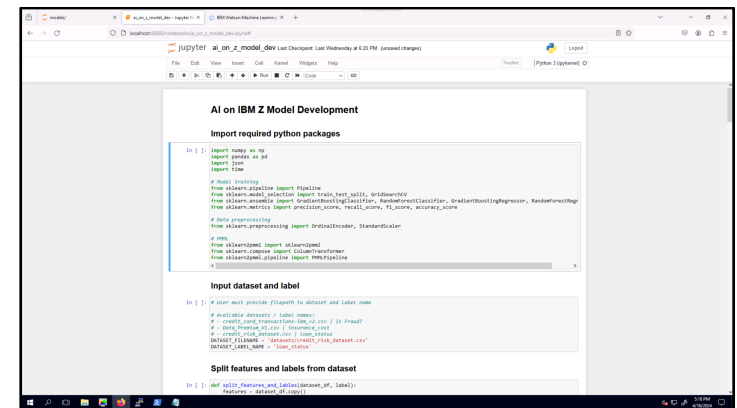
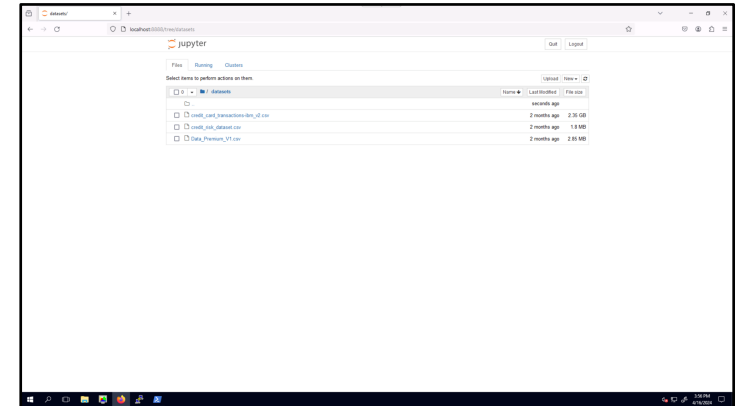
Provide data

1. Your input dataset (csv) in `datasets/` directory
2. Add input data to Jupyter notebook (`ai_on_z_model_dev.ipynb`)
 - a. Set `DATASET_FILENAME` to the path to your dataset
 - b. Set `DATASET_LABEL_NAME` to the name of the column you're predicting from the dataset

Model training

1. Step through and run all cells within Jupyter notebook (`ai_on_z_model_dev.ipynb`) within web browser

Note: This may take several minutes



1. AI model training

Access rapid AI on Linux
on Z development
environment

Provide data

Model training

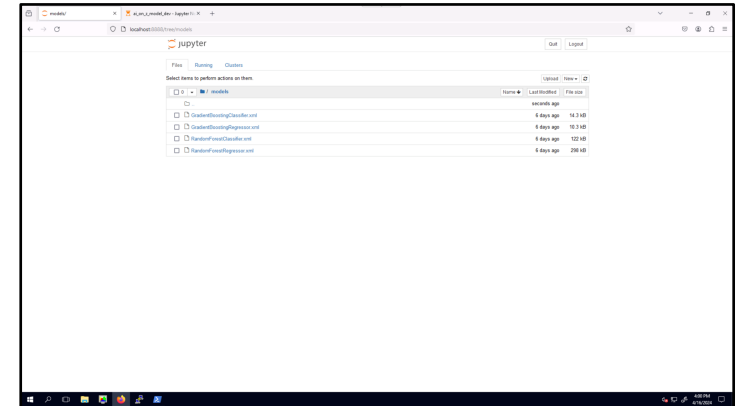
[Access trained AI model](#)

2. AI model deployment

Access trained AI model

1. Once training is complete, you can find your AI models within the `models/` directory (choose one for the following AI model deployment step)

3. AI model integration



✓ 1. AI model training

○ 2. AI model deployment

○ 3. AI model integration

☑ AI model training complete



Prerequisites

- Must have Docker or Podman installed

Step 2

AI model deployment

We will deploy our health insurance claims AI model using TIS. We can utilize the AI Toolkit to leverage TIS for model deployment. This deployed AI model can then be integrated into applications within the Linux on Z environment.

All sample code for this section is within

```
aionz-st-health-insurance-claims-tis/zST-model-deployment
```


Build Triton Inference Server

Integrate AI model into Triton Inference Server

Deploy Triton Inference Server

Run sample test

Build Triton Inference Server

1. Build podman image

```
podman build -t zst-tis .
```

Integrate AI model into Triton Inference Server

1. Add your model (.pmml file) to


```
aionz-st-health-insurance-claims-tis/zST-model-deployment/zST/models/rf_model/1
```

directory

2. Add your preprocessing .joblib file to

```
aionz-st-health-insurance-claims-tis/zST-model-deployment/zST/models/rf_model/1
```

directory



1. AI model training

2. AI model deployment

3. AI model integration

Build Triton Inference Server

Integrate AI model into Triton Inference Server

Deploy Triton Inference Server

Run sample test

Deploy Triton Inference Server

1. Create docker network

```
podman network create my-data-network
```

2. Run podman container

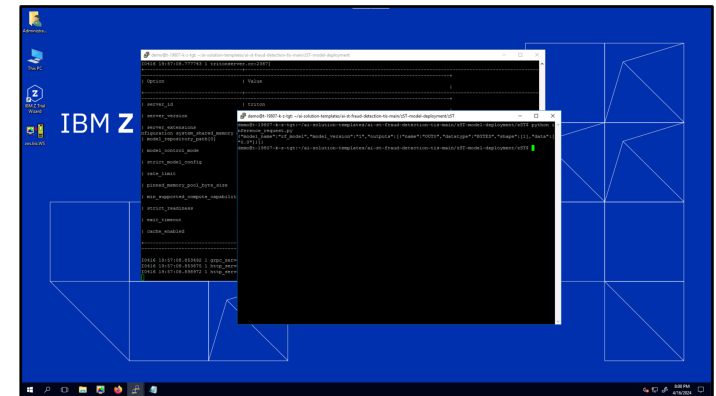
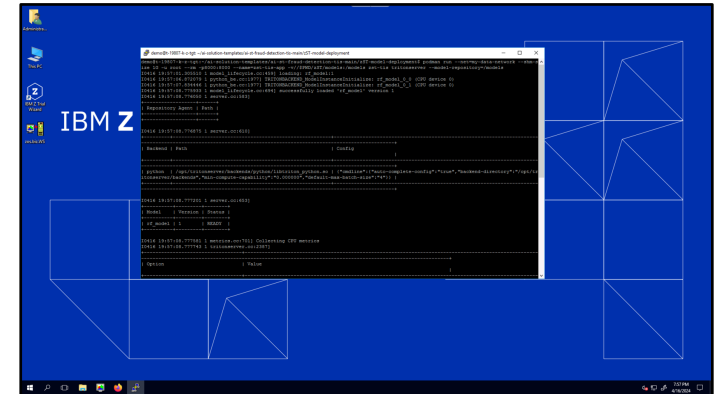
```
podman run --net=my-data-network --shm-size 1G  
-u root --rm -p8000:8000 --name=zst-tis-app -  
v//$PWD/zST/models:/models zst-tis-  
tritonserver --model-repository=/models
```

Run sample test

1. Run python script from terminal with ip/port of triton inference server (in new terminal)

```
cd aionz-st-health-insurance-claims-tis/zST-  
model-deployment/zST
```

```
python inference_request.py
```



✓ 1. AI model training

✓ 2. AI model deployment

○ 3. AI model integration

☑ AI model deployment complete



Prerequisites

- Must have Docker or Podman installed

Step 3

AI model integration

We can use our deployed TIS health insurance claims AI model and integrate it into different types of applications. The AI model can be analyzed and/or provide inferencing APIs using the sample AI on Linux on Z health insurance claims dashboard.

All sample code for this section is within

```
aionz-st-health-insurance-claims-tis/zST-model-  
integration-HIC
```

Configure sample application

Build sample application

Deploy sample application

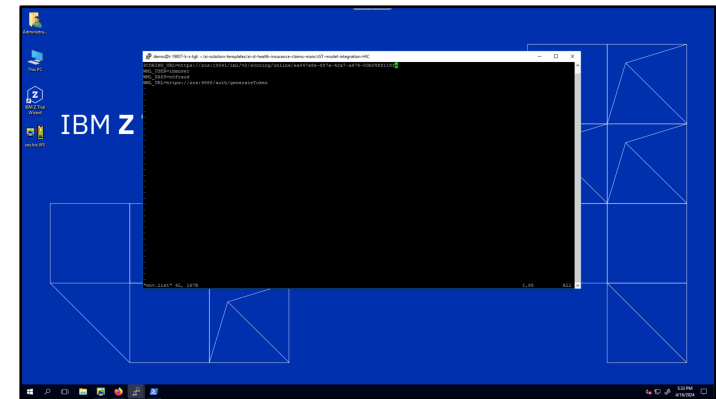
Access sample application

Configure sample application

1. Set the environment variables within

```
aionz-st-health-insurance-claims-tis/zST-model-integration-HIC/env.list
```

TIS_ENDPOINT (scoring endpoint for deployed AI model)



Build sample application

1. Run command in terminal

```
podman build -t health-insurance .
```

1. AI model training

Configure sample application

Build sample application

Deploy sample application

Access sample application

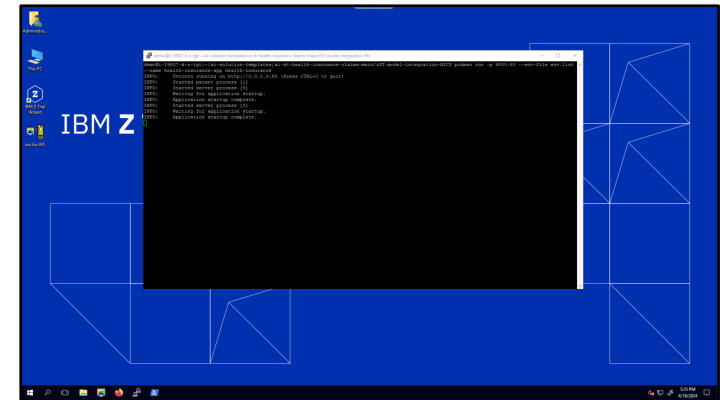
2. AI model deployment

Deploy sample application

1. Run command in terminal (e.g. port 9000)

```
podman run -p 9000:80 --env-file  
env.list --name health-insurance-app  
health-insurance
```

3. AI model integration

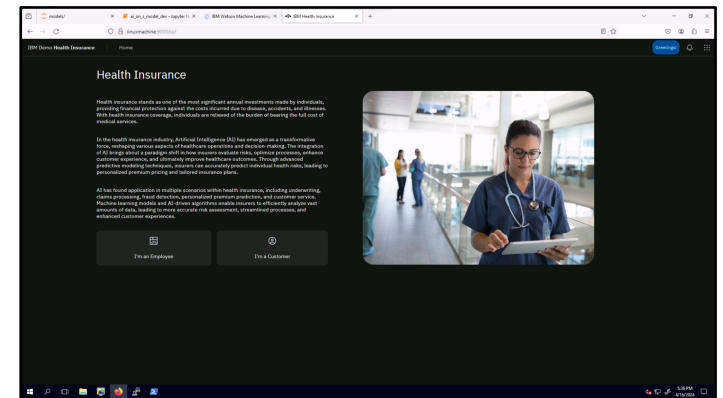


Access sample application

1. View the following URL in a web browser

<http://{ip address}:{port}/ui/>

ip address: IP of server you deployed application in
port: port you used with podman run



✓ 1. AI model training

✓ 2. AI model deployment

✓ 3. AI model integration

✓ AI model integration complete